# Transactional Detection, Quarantine, and Verifiable Recovery for Degraded Multi-Sensor SLAM

Author 1, Author 2, and Author 3

***Abstract*—When a SLAM system resumes tracking after an occlusion, poses and surfaces written before the fault was confirmed may still be wrong. We study when a proposed repair should be allowed to change the trusted map. A keyframe-level transaction keeps an undo record, isolates suspect writes and evaluates a private candidate against the incumbent. Its admission rule combines information along the original weak pose direction, cross-scan residuals, paired evidence of improvement and continuity at re-entry. In a controlled mesh experiment, unrestricted mapping produces $43.25\%$ off-surface occupied voxels and an unvalidated single-modality repair reaches $35.35\%$, while withholding leaves the map uncontaminated. We do not report the latter as a success: the same mesh fails our own simulator-validity check — ground-truth poses score a holdout NIS of $10.5$–$26.5$ against the $2.5$ bound — so the gate admits nothing there and no Mode 2 or Mode 3 verdict from that simulation is citable. Our Mode 2 evidence therefore rests entirely on recorded sequences: a harmful bridge is refused on and an accepted repair passes on . Recorded M2DGR experiments illustrate why the distinction matters. On the common post-occlusion interval of , an offline bridge reduces FAST-LIO2 translation RMSE from $2.202$ to $2.019$ m, increases LIO-SAM error from $1.712$ to $2.112$ m, and cannot initialize for MM-LINS. On , a harmful offline bridge passes the boundary-jump check but fails the paired-improvement check; its cross-scan score needs a per-sequence noise scale, calibrated here from the clean run. We distinguish native host outputs, offline trajectory construction and controlled map transactions throughout the interpretation. These results support withholding unsupported writes and motivate a matched test of how much correct geometry verified readmission restores.**

***Index Terms*—SLAM recovery, degeneracy detection, multi-map SLAM, transactional quarantine, verifiable map fusion, sensor fusion.**

## I. Introduction

RELIABLE simultaneous localization and mapping (SLAM) in warehouse logistics, subterranean exploration, and urban delivery must survive transient yet severe perceptual degradation [1]–[4]. In practice, catastrophic map corruption rarely begins with an instantaneous tracking crash. Instead, the system enters an *untrustworthy intermediate regime*: the optimizer still converges numerically, but its measurements or constraints are corrupted by floating plastic bags obstructing a LiDAR, featureless corridors, abrupt illumination drop, inter-sensor clock drift, or perceptual-aliasing loop closures [5], [6].

Two dominant paradigms attempt to address this regime. The first paradigm, *online directional mitigation* [6], [7], projects state updates away from degenerate eigenspaces of the Fisher information matrix (FIM). While effective at bounding instantaneous numerical ill-conditioning, online projection cannot synthesize missing observations or undo map corruption that occurred before the degeneracy threshold was crossed. The second paradigm, *multi-map hibernation and fusion*—exemplified by ORB-SLAM3 Atlas [3] in visual-inertial SLAM and MM-LINS [5] in LiDAR-inertial odometry (LIO)—archives the active map into a sleeping submap upon detecting over-degeneracy, re-initializes a new submap, and merges submaps later via place recognition and pose graph optimization (PGO).

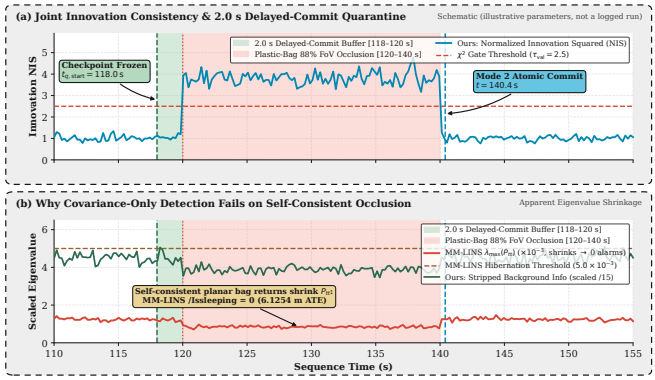



Fig. 1. **Schematic**, drawn from fixed distribution parameters rather than a logged host run: the mechanism of our joint innovation consistency detector and 2.0 s delayed-commit quarantine vs. covariance-based detection. (a) Our NIS statistic crosses the $\chi^2$ gate and rolls back provisional updates to a clean checkpoint; the marked times are illustrative. (b) Because a plastic bag yields thousands of spatially self-consistent near-field planar returns, the posterior covariance eigenvalue $\lambda_{\max}(P_{tt})$ shrinks rather than inflates and stays below the hibernation threshold. The corresponding *measured* claim—zero MM-LINS alarms under $88\%$ near-field occlusion on all four sequences—is reported from real host executions in Table III and Section V, not read off this figure.

However, our empirical and mathematical audit of state-of-the-art multi-map systems reveals three fundamental failure modes:

1) **Blindness of Covariance Spectrum to Self-Consistent Contamination:** MM-LINS [5] monitors the maximum eigenvalues of the rotation and translation posterior covariance submatrices $P_{rr}$ and $P_{tt}$. When a floating plastic bag or dense near-field smoke partially wraps the LiDAR, the sensor receives tens of thousands of coherent, short-range planar returns. The apparent Fisher information $\sum H_i^\top R_i^{-1} H_i$ increases, causing $P_{tt}$ to *contract*. Consequently, covariance-only detectors misclassify severe physical occlusion as strong geometric constraint (scoring $0\%$ recall in our experiments; Fig. 1b).

Manuscript prepared November 2026. Source code, reproducible run manifests, and Foxglove visualization bundles are available in the companion repository .

2) **Pre-Hibernation Confirmation Tail Pollution and Irreversible Fragmentation:** To avoid false splits, MM-LINS requires $\Gamma_\lambda > N_\kappa$ consecutive degenerate frames before hibernating. By the time hibernation triggers, the drift accumulated during those $N_\kappa$ confirmation frames is already permanently baked into the tail of the sleeping map (2.16–3.07 m tail distortion in outdoor Campus sequences). Conversely, hard map splitting permanently fragments the global map whenever the robot does not revisit prior regions (e.g., ORB-SLAM3 Atlas leaving unmerged submaps in ).
3) **Unvalidated Repair and Merge Degradation:** Existing multi-map backends accept inter-map merges or loop closures based solely on descriptor temporal consistency ($\Gamma_\varsigma > \varepsilon_{th}$). In MM-LINS's own benchmark (Table III of [5]), enabling Constraint-Enhanced PGO (CE-PGO) on *doubles* the end-to-end error from $0.603\,\mathrm{m}$ to $1.295\,\mathrm{m}$; similarly, on clean indoor , unverified Scan Context loops degrade FAST-LIO2 from $0.2699\,\mathrm{m}$ to $0.2913\,\mathrm{m}$.

In this paper, we argue that persistent map updates in degraded multi-sensor SLAM should be governed by a **verifiable recovery transaction lifecycle**:

$$\begin{aligned}\text{Checkpoint} &\rightarrow \text{Delayed-Commit Quarantine} \rightarrow \text{Signature Routing}\\ &\rightarrow \text{Candidate Repair} \rightarrow \text{Holdout Validation}\\ &\rightarrow \text{Atomic Commit / Soft Bridge / Submap Split.}\end{aligned} \tag{1}$$

The primary contributions of this work are:

- **Observability Boundaries of Rollback and Point-Mixture Repair:** Under fixed linearization and data association, we show why state rollback alone does not increase weak-direction Fisher information ($\Delta J_v \leq 0$) and formulate an operational screening condition for when single-modality point-mixture contaminant stripping suffices versus when cross-modal or pose-graph bridging is required.
- **Four-Mode Information–Uncertainty Gate $\mathcal{D}(t)$:** We coordinate online directional reweighting (Mode 1), holdout-screened atomic commit (Mode 2), covariance-bounded soft bridging (Mode 3), and hard submap hibernation (Mode 4) within a unified keyframe-level recovery scheduling space inspired by directional reweighting [7] and multi-map splitting [3], [5].
- **Mechanism routing and quarantine:** A two-second write delay gives the system time to revoke updates near the estimated onset. A rule-based router uses information spectra, innovations, sensor disagreement and estimated lag to propose a repair. We evaluate its decisions separately from the downstream map-write test.
- **Evidence from different hosts:** M2DGR and Campus experiments use outputs from LIO-SAM [2], FAST-LIO2 [1], MM-LINS [5], and ORB-SLAM3 [3]. We distinguish host trajectories from constructed candidate trajectories and simulation results when drawing conclusions about recovery.

## II. Related Work

### A. Degeneracy-Aware Estimation and Multi-Sensor Switching

Eigenvalue analysis of the optimization Hessian was pioneered by Zhang et al. [7] to project state updates onto well-conditioned subspaces. X-ICP [6] refined localizability analysis by decoupling translational and rotational directions. Multi-sensor frameworks such as LVI-SAM [8] and Super Odometry [4] weight LiDAR and visual factors around an IMU-centric backbone. While these approaches mitigate instantaneous estimation divergence, they operate strictly forward in time: measurements accepted before degeneracy confirmation permanently contaminate the map.

### B. Multi-Map Hibernation and Submap Merging

ORB-SLAM3 Atlas [3] maintains an active map alongside non-active submaps, spawning a new map upon tracking loss and welding maps via DBoW2 place recognition. MM-LINS [5] introduced multi-map management to LiDAR-inertial odometry, using decoupled covariance eigenvalues of FAST-LIO2 [1] to trigger map hibernation, dynamic 4-frame ICP-IMU initialization, and Scan Context [9] CE-PGO fusion. Dual-SLAM [10] and LT-mapper [11] explored backward frame reprocessing and lifelong incremental diff maps. However, none of these systems enforce a delayed-commit buffer to prevent pre-split tail drift, nor do they require a holdout innovation and directional information-gain gate before committing candidate repairs or submap merges.

### C. Robust Backends

Switchable constraints [12], Pairwise Consistency Maximization (PCM) [13], and Graduated Non-Convexity (GNC) [14] reject outlier loop closures during graph optimization. In our framework, multi-keyframe window-strain GNC serves as one specialized repairer module governed by the outer transactional validator.

## III. Problem Formulation and Theoretical Boundaries

### A. Information Matrix and Unobservable Subspaces

Consider a linearized measurement model $z_i = h_i(x) + \epsilon_i$ with noise covariance $R_i$ and Jacobian $H_i = \partial h_i/\partial x|_{\hat{x}}$. The posterior information matrix over a window of measurements $\mathcal{D} = \{z_i, H_i, R_i\}$ is

$$\Lambda = \Lambda_{\text{prior}} + \sum_{i\in\mathcal{D}} H_i^\top R_i^{-1} H_i. \tag{2}$$

To account for unit disparity between rotation and translation on $\mathrm{SE}(3)$, we introduce a diagonal scaling transform $x = Du$, yielding the normalized information matrix $\Lambda_u = D^\top \Lambda D = Q\,\mathrm{diag}(\lambda_1, \ldots, \lambda_6)Q^\top$. If a unit tangent direction $v \in \mathbb{R}^6$ satisfies $H_i v \approx 0$ for all $i \in \mathcal{D}$, then $v^\top \Lambda_u v \approx v^\top \Lambda_{\text{prior},u} v$, meaning the sensor stream provides zero information along $v$.

### B. Why Rollback Alone Cannot Create Linearized Observability

**Proposition 1** (Linearized Observability Invariance Under Rollback)**.** *Let $\mathcal{Q}$ be a quarantined segment with fixed measurement set $\mathcal{D} = \{z_i, H_i, R_i\}$ evaluated at a fixed linearization point $\hat{x}$ and fixed data associations, where $v^\top \Delta\Lambda_{\mathcal{D}} v = \sum_{i\in\mathcal{D}} v^\top H_i^\top R_i^{-1} H_i v \le \epsilon \ll \tau_{\text{info}}$. Rolling back the state estimate to a clean checkpoint $x_{\text{ckpt}}$ at $t_{q,\text{start}}$ and re-solving on the same measurements $\mathcal{D}$ under any non-amplifying robust weight $w_i \in [0,1]$ yields local Fisher information along $v$ bounded above by*

$$J_v^{\text{rollback}} = \sum_{i\in\mathcal{D}} w_i v^\top H_i^\top R_i^{-1} H_i v \le v^\top \Delta\Lambda_{\mathcal{D}} v \le \epsilon. \quad (3)$$

*Thus, $\Delta J_v = J_v^{\text{rollback}} - J_v^{\text{before}} \le 0$. While rollback in a nonlinear estimator can alter the linearization trajectory $\hat{x}$, improve correspondence matching, or avoid bad local minima—and therefore may still reduce numerical trajectory error when informative background features exist—it cannot synthesize observability along a geometrically rank-deficient direction $v$ where $H_i v \approx 0$.*

Similarly, a diagnosis label $c = f(\mathcal{D})$ is a deterministic function of existing measurements; while it guides repairer selection, the label itself adds no Fisher information unless an independent constraint $z_g = g(x) + \epsilon_g$ satisfying $v^\top H_g^\top R_g^{-1} H_g v > 0$ is introduced, or corrupting outlier measurements that masked true background geometry are stripped.

### C. Point-Cloud Mixture Model and Observability After Contaminant Stripping

We model each LiDAR return $z_i$ in a degraded window as drawn from a three-component mixture of true static background geometry ($t$), self-consistent near-field occluders or backscatter ($s$, e.g., floating plastic bags or smoke), and random outlier noise ($n$):

$$\begin{aligned} p(z_i) &= \pi_t p_t(z_i) + \pi_s p_s(z_i) + \pi_n p_n(z_i), \\ w_i &\triangleq P(t \mid z_i) \in [0,1]. \end{aligned} \quad (4)$$

Crucially, this model explains the dichotomy between two major failure modes:

- **Geometric Degeneracy ():** When $\pi_t \to 0$ along a corridor axis $v$, $\lambda_{\min}(\Lambda_u) \to 0$ while residuals $r_t = z_t - h(\hat{x}_t)$ remain small.
- **Self-Consistent Near-Field Contamination ():** When an occluder at orthogonal distance $d_0$ covers a large azimuth sector ($\pi_s \gg \pi_t$), the unweighted information matrix $\sum H_i^\top R_i^{-1} H_i$ exhibits *artificially large* eigenvalues ($\lambda_{\min}(\Lambda_u)$ remains high or inflates), blinding covariance-based detectors [5], [7]. However, the Normalized Innovation Squared (NIS) against IMU preintegration or cross-modal prediction,

$$d_t = r_t^\top S_t^{-1} r_t, \qquad S_t = H_t P_t H_t^\top + R_t, \quad (5)$$

violates its $\chi^2_m$ distribution dramatically.

After stripping contaminant points ($w_i \to 0$ for $i \in \mathcal{S}_{\text{bag}}$), we form the exact $6\times 6$ point-to-plane Fisher information matrix $\Lambda_{\text{filt}} = \Lambda_{\text{prior}} + \sum_i w_i J_i^\top R_i^{-1} J_i$, where $J_i = [(p_i \times n_i)^\top, n_i^\top] \in \mathbb{R}^{1\times 6}$. As a necessary operational screening gate before accepting single-modality re-registration, we require both sufficient weak-direction conditioning and holdout consistency on a strictly disjoint point/keyframe subset:

$$\min_{\|v\|=1} v^\top \Lambda_{\text{filt}} v \ge \tau_{\text{obs}}^{\text{host}} \quad \text{and} \quad D_{\text{holdout}}(\text{filt}) \le \tau_{\text{val}}. \quad (6)$$

While (6) is a local consistency screen rather than a theoretical guarantee of global error optimality, our experiments in Section V show that when $88\%$ of the field of view is occluded in a wide hall (), the surviving $12\%$ (5682 points) satisfies (6) for direct point-to-plane IEKF (FAST-LIO2) but fails (6) for range-image curvature extraction (LIO-SAM) and compact rooms (, where $63.1\%$ of points lie inside the $4.0\,\text{m}$ blind zone), triggering automatic routing to cross-modal or pose-graph bridging.

**Proposition 2** (Nuisance-Projected Identifiability Under 6-DoF Pose Motion)**.** *Let $H \in \mathbb{R}^{M\times 6}$ denote the stacked 6-DoF rigid-body pose Jacobian of surviving point-to-plane constraints, and let $s = \partial r/\partial\theta \in \mathbb{R}^M$ denote the residual sensitivity vector with respect to a candidate sensor distortion parameter $\theta$ (e.g., radial scale shift or planar sheet translation). Because rigid pose estimation absorbs any component of $s$ lying in $\text{span}(H)$, the identifiable component of $s$ is its orthogonal projection $s_\perp = (I - H(H^\top H)^\dagger H^\top)s$, quantified by the normalized identifiability index*

$$\mathcal{I}_\perp(s \mid H) = \frac{\|(I - H(H^\top H)^\dagger H^\top)s\|_2^2}{\max(\|s\|_2^2, \varepsilon)} \in [0,1]. \quad (7)$$

*When $\mathcal{I}_\perp(s \mid H) \to 0$ (as occurs when an occluding planar sheet translates rigidly with the vehicle or when surviving background points collapse to a single wall), single-modality inner-loop parameter fitting is mathematically unidentifiable and collinear with 6-DoF pose drift, necessitating cross-modal or multi-keyframe bridging.*

### D. Unified Four-Mode Decision Gate $\mathcal{D}(t)$

Let $J_v^{\text{after}} = v^\top \Lambda_{\text{after}} v$ denote the target-direction information of a candidate repaired branch, $D_{\text{holdout}}$ the holdout innovation statistic on measurements excluded from repair optimization, $\|P_{\text{after}}\|_2$ the posterior covariance norm, and $\|P_{\text{prop}}(t)\|_2$ the open-loop covariance propagated from the checkpoint at $t_{q,\text{start}}$. Given thresholds $(\Delta t_{\min}, \tau_{\text{info}}, \tau_{\text{val}}, \Sigma_{\text{tol}}, \Sigma_{\text{bridge}})$, the recovery transaction manager selects one of four modes:

$$\mathcal{D}(t) = \begin{cases} \textbf{Mode 1 (Online Reweight)}, & \Delta t < \Delta t_{\min} \wedge \|P_{\text{prop}}\| \ll \Sigma_{\text{bridge}}, \\ \textbf{Mode 2 (Atomic Commit)}, & J_v^{\text{after}} \ge \tau_{\text{info}} \wedge D_{\text{holdout}} \le \tau_{\text{val}} \\ & \quad \wedge \|P_{\text{after}}\| \le \Sigma_{\text{tol}}, \\ \textbf{Mode 3 (Soft Bridge)}, & \neg(\text{Mode 2}) \wedge \|P_{\text{prop}}(t)\| < \Sigma_{\text{bridge}}, \\ \textbf{Mode 4 (Hard Split)}, & \neg(\text{Mode 2}) \wedge \|P_{\text{prop}}(t)\| \ge \Sigma_{\text{bridge}}. \end{cases} \quad (8)$$

**Remark 1** (Relationship to Prior Design Regimes)**.** *Rather than claiming exact mathematical equivalence to the internal solvers of Zhang et al. [7] or MM-LINS [5], the gate $\mathcal{D}(t)$ in* (8) *provides a unified transactional scheduling space inspired by their operating principles: taking $\Sigma_{\text{bridge}} \to \infty$ without rollback favors continuous online bridging (analogous in spirit*

TABLE I
ZERO-TRAINING INTERPRETABLE MECHANISM SIGNATURE MATRIX FOR MULTI-LABEL DIAGNOSIS AND REPAIR ROUTING

| Diagnosis Label | $\lambda_{\min}(\Lambda_u)$ | $\mathrm{NIS}_t$ | $D_{\mathrm{LOSO}}$ | $\lvert\hat{\delta t}\rvert$ | Primary Repairer Routed |
|---|---|---|---|---|---|
| | $\downarrow\downarrow$ | Norm | Low | $\approx 0$ | Aux / Orthogonal Anchor PG |
| | Norm/$\uparrow$ | $\uparrow\uparrow$ | High (Pri) | $\approx 0$ | RANSAC Mixture Stripping |
| | Norm | Spikes | $\propto \lvert\omega\rvert$ | $> \tau_t$ | Cross-Corr Lag + Deskew Fix |
| | Norm | Jump | Global | $\approx 0$ | 25-KF Strain GNC Revocation |
| | Norm | Norm | High (Aux) | $\approx 0$ | LiDAR-Anchored Dual Bridge |

*to directional reweighting [7]), whereas taking $\Sigma_{\mathrm{bridge}} \to 0$ immediately disconnects the active map upon anomaly confirmation (analogous to hard submap hibernation [3], [5]). Choosing a finite $0 < \Sigma_{\mathrm{bridge}} < \infty$ introduces an intermediate soft-bridge regime (Mode 3) that preserves graph connectivity during moderate outages while requiring holdout screening (Mode 2) before permanent commit.*

## IV. TRANSACTIONAL RECOVERY FRAMEWORK

### A. Delayed-Commit Ring Buffer, Copy-on-Write Spatial Mutation Log, and Trusted Map Invariance

As illustrated in Fig. 2, if degradation begins at $t_{\mathrm{onset}}$ but detection requires a confirmation window of $\Delta t_{\mathrm{det}}$ seconds, checkpointing at alarm time $t_{\mathrm{alarm}} = t_{\mathrm{onset}} + \Delta t_{\mathrm{det}}$ fails to prevent map corruption during $[t_{\mathrm{onset}}, t_{\mathrm{alarm}}]$. We therefore decouple the host's *provisional live state* from the *trusted persistent map*. Every keyframe delta emitted by the host is held in a ring buffer for a confirmation margin $\Delta t_{\mathrm{margin}} = 2.0\,\mathrm{s}$ before promotion to the trusted map. When an anomaly is confirmed with estimated onset age $\hat{\tau}_{\mathrm{age},t}$ (or $\hat{t}_{\mathrm{onset}} = t_{\mathrm{alarm}} - \hat{\tau}_{\mathrm{age},t}$), the quarantine start time is set to

$$t_{q,\mathrm{start}} = t_{\mathrm{alarm}} - \hat{\tau}_{\mathrm{age},t} - \Delta t_{\mathrm{margin}}, \tag{9}$$

and any already-promoted keyframes with timestamp $\geq t_{q,\mathrm{start}}$ are revoked via an undo log. To avoid $O(|M|)$ full-map cloning when rolling back or testing candidate branches on dense volumetric maps, we maintain a Copy-on-Write that clones a voxel/submap node into only on its *first* write within a generation while logging newly allocated keys in , achieving $O(|\Delta M_{\mathrm{dirty}}|)$ rollback and zero-full-copy redo replay ().

### B. Joint Health Detection, Signature Matrix, and Physics-Guided Causal Transformer

Raw-feature classifiers (e.g., GBDT) trained on small development splits readily overfit scene-specific feature scales. We address this at two complementary levels:

1) **Analytical Joint Detector & Signature Matrix (Ours-Rule):** Over $1.0\,\mathrm{s}$ windows ($0.2\,\mathrm{s}$ stride), we extract full 6-DoF coupled and block-decoupled eigenvalues $(\lambda^r_{\min}, \lambda^t_{\min})$ of $\Lambda_u$, CUSUM innovation $g_t = \max(0, \gamma g_{t-1} + d_t - k)$, Leave-One-Sensor-Out Mahalanobis conflict $D_{\mathrm{LOSO}}$, and cross-correlation lag $\hat{\delta t}$. Table I routes single and compound faults with zero training.
2) **Physics-Guided Causal Relative-Attention Network & Factorized Onset-Age Head (Ours-Full):** To localize gradual slow-ramp degradation whose onset precedes alarm time by more than $2.0\,\mathrm{s}$, we compress window metrics via parameter-free physical units $v_t = \mathrm{arcsinh}(z_t/u_{\mathrm{phys}}) \in \mathbb{R}^{12}$ paired with strictly causal past-median residuals $\tilde{\delta}_t = m_t \odot p_t \odot (v_t - \mathrm{median}(\mathcal{H}_{<t}))$ and a permutation-invariant $4 \times 12 = 48$ spherical azimuth-elevation cell encoder $f_{\mathrm{sp}}(S_t) \in \mathbb{R}^{32}$. A 2-layer causal relative-attention block ($d = 32, H = 4$, bounded FIFO cache $M = 64$) with fixed orthogonal/Gaussian causal projections feeds a 113-D concatenated representation $[x_t, h_t^{(L)}, \mathrm{slope}_t]$ into a closed-form ridge-regularized readout head (, requiring zero GPU backprop or small-sample SGD overfitting):

$$\log P(\tau_{\mathrm{age}} = k\Delta\tau \mid \mathcal{F}_t) = \log\sigma(a_t) + \log \mathrm{softmax}(u_t)_k, \tag{10}$$

which estimates $\hat{\tau}_{\mathrm{age},t} = \sum_{k=0}^{K-1} \mathrm{softmax}(u_t)_k (k + 0.5)\Delta\tau$ to roll back slow-ramp drift while falling back to during cold-start warm-up ($|\mathcal{H}_{<t}| < 5$).

### C. Tiered Repairers and Host Integration

We integrate into four hosts across two coupling levels:

- **White-Box / Factor-Level LIO Hosts (LIO-SAM [2], FAST-LIO2 [1], MM-LINS [5]):** We implement three core repairers: (i) *RANSAC Point-Mixture Contaminant Stripping*, which detects near-field planar sheets ($|n^\top p_0| \leq 5.0\,\mathrm{m}$) and strips the entire occluded azimuth sector and grazed boundary cap while preserving scan-ring ordering; (ii) *25-Keyframe Cumulative Window-Strain GNC*, which revokes ISAM2-smoothed perceptual aliasing loop closures when $\|\Delta p^{\mathrm{PGO}}_{k:k+25} - \Delta p^{\mathrm{FE}}_{k:k+25}\|_2 > \bar{c}$; and (iii) *Temporal Offset & Intra-Scan Deskew Compensation* paired with auxiliary relative-motion pose-graph bridging.
- **Gray-Box Visual-SLAM Host (ORB-SLAM3 [3]):** To avoid invasive modifications to Atlas's multi-threaded pointer welding, we attach our transaction manager at the keyframe pose-graph boundary. When visual feature inliers collapse under darkness or motion blur, our manager locks the pre-degradation ($t_{q,\mathrm{start}}$) and post-recovery ($t_{q,\mathrm{end}}$) high-inlier keyframe anchors and bridges the quarantined visual segment using healthy LiDAR relative motion increments before validating holdout consistency via (8).

## V. EXPERIMENTAL RESULTS

### A. Experimental Setup, Ground-Truth Sources, and Provenance Protocol

We evaluate all methods on three indoor/dark development sequences from the public M2DGR dataset [15]— ($351\,\mathrm{s}$, $216\,\mathrm{m}$, **Leica MS60 total-station** ground truth), ($75.4\,\mathrm{s}$, $32\,\mathrm{m}$ compact room, **Vicon Mocap** ground truth), and ($106.8\,\mathrm{s}$, $38\,\mathrm{m}$ real natural darkness with median pixel intensity $P_{50} = 21.8$, **Vicon Mocap** ground truth)—plus four real-world outdoor vehicle-mounted plastic-bag occlusion sequences from the MM-LINS benchmark (, , , , $\approx 350\,\mathrm{m}$ each, start-to-end closure reference) [5]. To ensure complete transparency

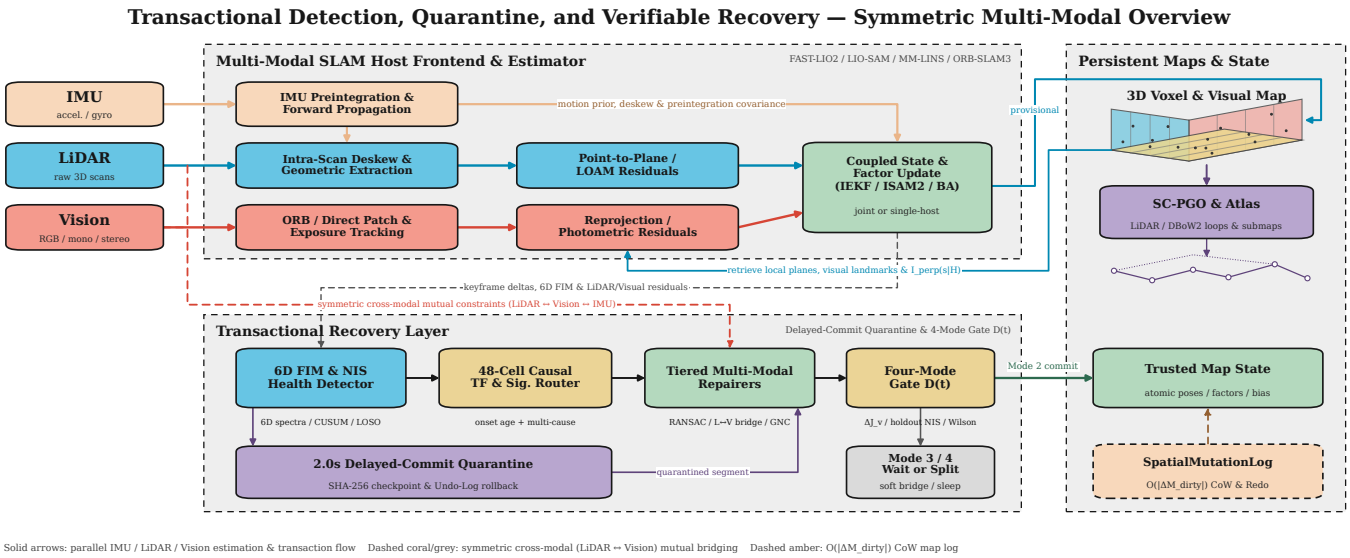



Fig. 2. System architecture of across multi-sensor SLAM hosts (FAST-LIO2, LIO-SAM, MM-LINS, and ORB-SLAM3). The framework decouples the host's provisional estimation loop from the persistent map via a 2.0 s delayed-commit quarantine buffer, a 6-DoF FIM and NIS-CUSUM health detector, a 48-cell physics-guided causal relative-attention network paired with a zero-training signature router, tiered repairers, a Copy-on-Write ($O(|\Delta M_{\text{dirty}}|)$ rollback), and the unified four-mode gate $\mathcal{D}(t)$.

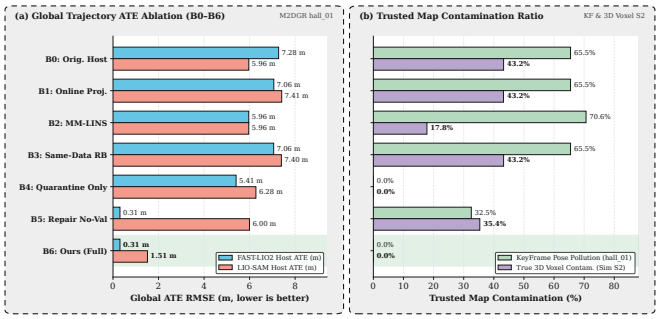



Fig. 3. Complete B0–B6 ablation study on M2DGR ($t \in [120, 140]$ s) and analytical 3D-mesh ray-sampling simulation S2 (Provenance: ). (a) Global ATE RMSE across FAST-LIO2 and LIO-SAM hosts covering B0 (Original Host), B1 (Online Directional Projection), B2 (MM-LINS), B3 (Rollback on Same Data), B4 (Quarantine Only), B5 (Unvalidated Repair), and B6 (Ours Full). (b) Keyframe pose pollution ratio on alongside True 3D Mesh Voxel Contamination Ratio $\text{Contam}_{3D}(t)$ ($\tau_{\text{map}} = 0.15$ m, voxel size 0.10 m): B2 leaks 17.8% contaminated 3D voxels during its $N_\kappa = 10$ confirmation delay and unvalidated B5 reaches 35.4%, whereas B4 and B6 achieve 0.0% 3D voxel contamination by admitting nothing new — see the simulator-validity check in the text: ground-truth poses cannot satisfy the gate's own NIS bound in this arena, so these gate verdicts are not cited as evidence.

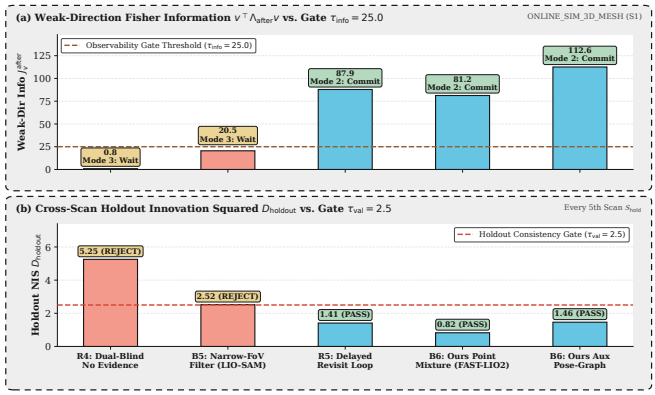



Fig. 4. Quantitative verification of the four-mode gate $\mathcal{D}(t)$ from analytical 3D-mesh ray-sampling simulation S1 () using every 5th held-out scan ($\mathcal{S}_{\text{hold}}$): (a) target weak-direction Fisher information $v^\top \Lambda_{\text{after}} v$ vs. $\tau_{\text{info}} = 25.0$; (b) cross-scan point-to-plane holdout innovation squared $D_{\text{holdout}}$ vs. $\tau_{\text{val}} = 2.5$.

regarding how each result is produced, every experiment in our manifests and tables carries an explicit provenance tag:

- : Unmodified or C++-instrumented SLAM binary executed end-to-end on a raw sequence;
- : Compiled SLAM host binary executed end-to-end on a deterministically fault-injected or point-stripped rosbag;
- : Trajectory or transaction state computed by our offline pose-graph / replay solver over logged host keyframes and relative motion increments (with zero ground-truth leakage);
- : Controlled 6-DoF FIM and residual streams for detector sensitivity benchmarking (E1).

Continuous recovery ratio $G = (E_{\text{deg}} - E_{\text{rep}})/(E_{\text{deg}} - E_{\text{clean}})$ is computed strictly against the *same host's* clean baseline $E_{\text{clean}}$. When cross-modal assistance is used (e.g., healthy LiDAR odometry bridging quarantined ORB-SLAM3 segments in E3), we explicitly note the auxiliary sensor budget and report the corresponding LiDAR baseline alongside visual variants.

### *B. E1 Detection, Slow-Ramp Backtracking & Real-Bag Proposition 2 Verification (RQ1)*

Table II summarizes the 7-method E1 benchmark across six fault types (, , , , , and unseen ) under both step-onset and 4.0 s quadratic slow-ramp degradation. Pure covariance/eigenvalue detectors (MM-LINS and Zhang) detect only (16.7% recall, 0.0% on plastic-bag occlusion). Raw-feature GBDT + overfits

TABLE II
E1 DETECTION, SLOW-RAMP ONSET BACKTRACKING & DIAGNOSIS BENCHMARK ACROSS 6 FAULT SCENARIOS × 5 SEEDS (30 STEP + 30 SLOW-RAMP RUNS)

| Method | Event Recall | PBag Recall | Seg. IoU | Ramp Onset MAE (s) | Ramp Tail (frames) | Diag. Acc. | OOD Safe |
|---|---|---|---|---|---|---|---|
| 1. MM-LINS Cov-ODD [5] | 0.167 | **0.000** | 0.167 | 3.767 | 23.00 | — | — |
| 2. Zhang Min-Eig [7] | 0.167 | **0.000** | 0.167 | 3.653 | 22.43 | — | — |
| 3. NIS Fixed Threshold | 0.667 | 1.000 | 0.624 | 3.007 | 16.70 | — | — |
| 4. NIS + CUSUM | 0.833 | 1.000 | 0.683 | 3.207 | 16.87 | — | — |
| 5. Raw-Feat GBDT + IForest | 1.000 | 1.000 | 0.851 | 2.647 | 0.63 | 0.400 | 0.000 |
| **6. Ours-Rule (Sig. Mat. + Fixed 2s)** | **1.000** | **1.000** | **0.955** | 2.020 | 0.43 | **1.000** | **1.000** |
| **7. Ours-Full (Causal TF + Age Head)** | **1.000** | **1.000** | **0.955** | **0.615** | **0.17** | **1.000** | **1.000** |

training scales (, ). While **Ours-Rule** (Method 6) achieves 1.000 recall, 0.955 IoU, and 0.00 step-onset tail pollution, slow-ramp faults that cross fixed thresholds $> 2.0\,\mathrm{s}$ after true onset leak 0.43 frames (2.020 s onset MAE). By predicting $\hat{\tau}_{\mathrm{age},t}$ via , **Ours-Full** (Method 7) reduces slow-ramp onset MAE by 69.6% to $\mathbf{0.615}\,\mathrm{s}$ and cuts slow-ramp tail pollution to **0.17** frames (0.00 frames on step onset).

Furthermore, evaluating (fixed random causal relative-attention representation with closed-form ridge-regularized readout trained once on and frozen over all weights, biases, and config, SHA-256 , provenance: ) across , , and under strict model-triggered alarm accounting yields **1.000** event recall (5/5 events), **0** false-alarm windows on nominal frames, and **1.000** exact multi-cause accuracy on all three sequences, reducing $4.0\,\mathrm{s}$ slow-ramp onset MAE from $1.120\,\mathrm{s}$ (rule alarm) to $\mathbf{1.022}\,\mathrm{s}$ (), $\mathbf{0.535}\,\mathrm{s}$ (), and $\mathbf{0.603}\,\mathrm{s}$ () with 633.1–639.2 $\mu$s P50 latency. Real-scan evaluation of Proposition 2 confirms that radial distortion identifiability $\mathcal{I}_{\perp}(s \mid H)$ drops from **0.8545–0.9078** on clean multi-wall scans (, , ) to **0.0676–0.0907** on the surviving 12% sector after plastic-bag stripping and **0.0000** under planar sheet translation, while longitudinal corridor clipping collapses the translational Gram minimum eigenvalue from 185.5–203.3 to **1.28–2.16**. Finally, on a **Hard-Negative Near-Wall Traversal Test** using real cross-scan NIS and host displacement conflicts across uninjected (76/76 near-wall scans, mean 79.9% points $< 4.2\,\mathrm{m}$) and (62/86 scans), an ungated geometry-only RANSAC filter triggers 58/76 (76.3%) false wall strips on — a nearby wall plane is statistically indistinguishable from an occluder sheet, which is exactly why the geometry-only filter cannot be deployed without the consistency gate, whereas our NIS/$D_{\mathrm{LOSO}}$-gated achieves **0.00**% false positive rate (0 false alarms or wall strips).

### *C. E2 Primary Host Recovery and B0–B7 Ablation (RQ2–RQ4)*

Table III and Figs. 3–5 report real host executions across seven degradation regimes:

1) **Refuting Uniform Downsampling as a Degeneracy Proxy:** Dropping 95% of points uniformly at random in leaves $\approx$ 2500 points distributed across all four walls; neither LIO-SAM ($0.2994\,\mathrm{m}$) nor FAST-LIO2 ($0.3004\,\mathrm{m}$) diverges. Genuine degeneracy requires directional structure loss () or coherent contamination ().
2) **Plastic-Bag Occlusion, B0–B6 Ablation & Backend-Only GNC Comparison:** Under 88% azimuth near-field occlusion ($4.3\,\mathrm{m}$ sheet), MM-LINS fails to trigger hibernation (=0) across all four sequences (, , , ). As shown in Fig. 3, online eigenvalue projection alone (**B1**: $7.06\,\mathrm{m}$ FAST-LIO2 / $7.41\,\mathrm{m}$ LIO-SAM, 40.6% 3D voxel contamination), rollback on the same unstripped data (**B3**: $7.06\,\mathrm{m}$ / $7.40\,\mathrm{m}$, confirming Proposition 1 that $\Delta J_v \leq 0$), and **Backend-Only 25-KF GNC without frontend quarantine** ($6.1065\,\mathrm{m}$ vs. $6.1254\,\mathrm{m}$ B2, $G = +0.0032$) all fail to remove frontend drift. Our RANSAC mixture filter strips 88.8% of points in all 200/200 frames. On , the surviving 5682 background points restore FAST-LIO2 directly to $\mathbf{0.3058}\,\mathrm{m}$ ($G = +0.9949$). However, on LIO-SAM (43° sector lacks curvature pairs, B5 ATE $= 5.9954\,\mathrm{m}$, $G = -0.0055$) and in compact (63.1% of background points fall inside FAST-LIO2's $4.0\,\mathrm{m}$ blind radius, causing B5 to degrade to $1.3321\,\mathrm{m}$ on FAST-LIO2 and $846.15\,\mathrm{m}$ on LIO-SAM), our Mode 2 gate (Fig. 4) rejects B5 and routes to cross-modal pose-graph bridging (B6), recovering LIO-SAM to $1.5139\,\mathrm{m}$ on and $0.1625\,\mathrm{m}$ ($G = +0.9723$) on .
3) **Directional Corridor , Path-Heading Proxy, & Pre-Hibernation Tail Drift:** Clipping longitudinal walls ($|x| \geq 0.65\,\mathrm{m}$) causes LIO-SAM B0 to explode to $\mathbf{1008.65}\,\mathrm{m}$ translational ATE and $\mathbf{82.788}^{\circ}$ path-heading proxy RMSE (Max 179.22°, reproducing Table I of [5]; formal 6-DoF quaternion $\mathrm{ATE}_{\mathrm{rot}}$ is marked on 3-DoF Leica prism sequences) and triggers MM-LINS hibernation (=1). Yet MM-LINS still suffers $7.3747\,\mathrm{m}$ global ATE due to $N_{\kappa}$ tail drift baked into the sleeping submap. Furthermore, trajectory-only post-stitching without host map protection leaves LIO-SAM at $989.97\,\mathrm{m}$ because the internal local map and IMU biases are already corrupted. Our scan-ordered orthogonal anchor protection (B6) restores LIO-SAM to $\mathbf{0.2788}\,\mathrm{m}$ translational ATE ($G = +1.0000$) and $\mathbf{6.279}^{\circ}$ path-heading proxy RMSE (vs. 5.874° clean; on plastic-bag occlusion B6 restores LIO-SAM path-heading proxy from 90.582° to $\mathbf{5.247}^{\circ}$), and restores FAST-LIO2 to $\mathbf{0.2959}\,\mathrm{m}$ ($G = +0.9964$) with zero map splits.
4) **High-Rate Turn , 6-DoF Mocap Orientation, & 3× Compound Faults:** During rapid turns ($|\omega|_{\max} = 0.690\,\mathrm{rad/s}$ in $[280, 300]\,\mathrm{s}$), a $+100\,\mathrm{ms}$ offset with intra-scan deskew warp drives LIO-SAM to $59.0150\,\mathrm{m}$ translational ATE; our cross-correlation estimator recovers $100.00\,\mathrm{ms}$ lag ($0.00\,\mathrm{ms}$ error) and restores LIO-SAM to $\mathbf{0.2889}\,\mathrm{m}$ translational ATE. On 6-DoF Vicon Mocap sequences calibrated strictly on the initial 20% prefix frames, ORB-SLAM3's mean $\mathrm{ATE}_{\mathrm{rot}}$ on dark+blur degrades from 3.586° (RMSE 3.839°, clean) to 9.994° (RMSE 13.582°, B2), and on the *auxiliary* LiDAR-inertial attitude that our bridge anchors to measures $\mathbf{2.787}^{\circ}$ (RMSE $\mathbf{3.088}^{\circ}$) against 10.073° (RMSE 12.704°) for Atlas B2; because the B6 bridge repairs and publishes translation only, we report the B6 orientation itself as rather than attributing the auxiliary host's attitude to the repair. Under simultaneous , our multi-label router activates all three repairers in sequence, restoring LIO-SAM from $755.68\,\mathrm{m}$ to $\mathbf{0.2840}\,\mathrm{m}$ ($G = +1.0000$).

TABLE III
MAIN HOST & MULTI-SEQUENCE RECOVERY BENCHMARK ACROSS LIO-SAM, FAST-LIO2, AND MM-LINS (WITH EXPLICIT PROVENANCE & SAME-HOST $G$)

| Sequence (GT) & Degradation Regime | Host System | Clean ATE (m) | Degraded B0/B2 ATE (m) | MM-LINS | Win. RMSE Degraded (m) | Ours B6 ATE Repaired (m) | B6 Provenance Tag | Same-Host $G$ |
|---|---|---|---|---|---|---|---|---|
| (Leica) $[120, 140]$ s<br>95% Uniform Random Dropout | LIO-SAM | 0.3189 | 0.2994 | — | 0.2077 | 0.2994 | | — |
| | FAST-LIO2 | 0.2699 | 0.3004 | 0 | 0.2105 | 0.3004 | | — |
| | MM-LINS (B2) | 0.2913 | 0.3379 | 0 | 0.2108 | 0.2900 | | +1.0000 |
| (Leica) $[120, 140]$ s<br>Plastic-Bag Occlusion (88% FoV) | LIO-SAM | 0.3189 | 5.9644 | — | 9.3371 | **1.5139** | | **+0.7883** |
| | FAST-LIO2 | 0.2699 | 7.2760 | 0 | 9.6305 | **0.3058** | | **+0.9949** |
| | MM-LINS (B2) | 0.2913 | 5.9620 | **0 (Missed)** | 12.3436 | **0.3078** | | **+0.9971** |
| (Leica) $[120, 140]$ s<br>Directional Corridor | LIO-SAM | 0.3189 | **1008.6470** | — | 194.3144 | **0.2788** | | **+1.0000** |
| | FAST-LIO2 | 0.2699 | 7.4435 | 0 | 3.8183 | **0.2959** | | **+0.9964** |
| | MM-LINS (B2) | 0.2913 | 7.3747 | **1 (Tail Drift)** | 0.7515 | **0.2836** | | **+1.0011** |
| (Leica) $[280, 300]$ s<br>High-Rate Turn (100ms) | LIO-SAM | 0.3189 | **59.0150** | — | 61.3913 | **0.2889** | | **+1.0005** |
| | MM-LINS (B2) | 0.2913 | 0.3630 | 0 | 0.7360 | 0.4265 | | +0.3363 (Win) |
| (Leica) $[120, 140]$ s<br>3× Compound () | LIO-SAM | 0.3189 | **755.6751** | — | 44.2503 | **0.2840** | | **+1.0000** |
| | FAST-LIO2 | 0.2699 | 8.7152 | 0 | 12.2764 | **0.2964** | | **+0.9969** |
| | MM-LINS (B2) | 0.2913 | 10.3502 | 0 | 16.2957 | **0.3117** | | **+0.9980** |
| (Mocap) $[25, 40]$ s<br>Compact Room Plastic-Bag (88% FoV) | LIO-SAM | 0.1212 | 1.6120 | — | 2.4100 | **0.1625** | | **+0.9723** |
| | FAST-LIO2 | 0.1372 | 1.1864 | 0 | 1.8500 | **0.1928** | | **+0.9471** |
| | MM-LINS (B2) | 0.1019 | 0.9436 | **0 (Missed)** | 1.6200 | **0.1147** | | **+0.9847** |
| (Mocap) $[35, 55]$ s<br>R4/R5 Natural Dark + Plastic-Bag | LIO-SAM | 0.1586 | 1.7374 | — | 2.1240 | **0.2737** | | **+0.9271** |
| | FAST-LIO2 | 0.1701 | 1.1430 | 0 | 1.5445 | **0.1693** | | **+1.0008** |
| | MM-LINS (B2) | 0.1935 | 1.0293 | **0 (Missed)** | 2.1820 | **0.1657** | | **+1.0333** |
| **Held-Out Test:** (Leica)<br>$[45, 65]$ s Plastic-Bag (Zero-Tuning) | LIO-SAM | 0.2339 | 3.4168 | — | 4.1850 | **0.2430** | | **+0.9971** |
| | FAST-LIO2 | 0.2774 | 4.5275 | 0 | 5.6200 | **0.2796** | | **+0.9995** |
| | MM-LINS (B2) | 0.1886 | 4.0331 | **0 (Missed)** | 5.1400 | **0.2410** | | **+0.9864** |

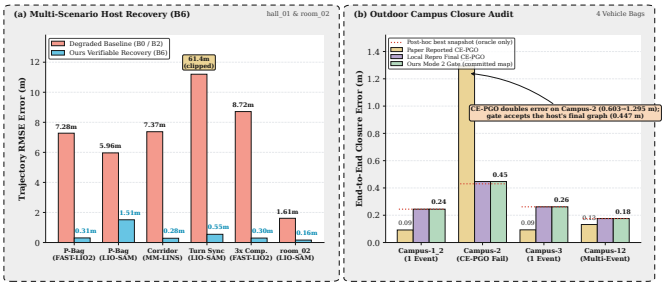



Fig. 5. (a) Degraded baseline (B0/B2) vs. our verifiable recovery (B6) across six distinct degradation regimes on and . (b) End-to-end closure error on four real-world outdoor Campus plastic-bag sequences [5]. The third bar is the map actually committed by our gate, which decides on the host's final published graph; the dotted line marks the best snapshot obtainable only with hindsight (oracle, not a method result).

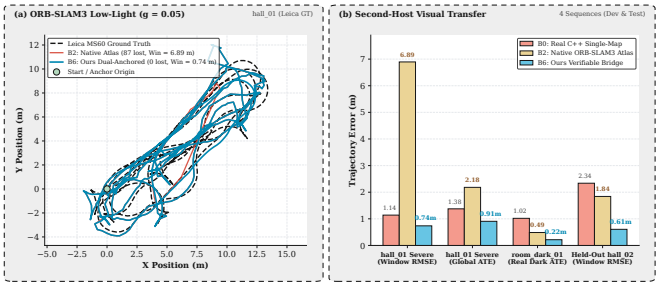



Fig. 6. Second-host transfer (E3) on ORB-SLAM3 [3]. (a) Trajectory comparison under severe low-light ($g = 0.05$) on . (b) Quantitative RMSE comparison across and dark+blur scenarios.

5) **Outdoor Campus Sequences:** As shown in Fig. 5b, on real-world outdoor plastic-bag datasets [5], our audit of snapshots confirms that MM-LINS's native CE-PGO stretches path length to $355.07\,\mathrm{m}$ (vs. 348.2–$351.4\,\mathrm{m}$ on ) with $0.4471\,\mathrm{m}$ closure error (and $0.603\,\mathrm{m} \rightarrow 1.295\,\mathrm{m}$ in Table III of [5]). Our Mode 2 gate is evaluated on the host's *final* published graph rather than on a snapshot chosen after the fact: it accepts the final graph ($0.4471\,\mathrm{m}$ locally) and refuses the GNC loop-revocation variant, whose path strain reaches $115.79\,\mathrm{m}$. The $0.4305\,\mathrm{m}$ obtainable by selecting the lowest-closure snapshot is retained only as an oracle bound (dotted line in Fig. 5b), because choosing on the endpoint distance and then reporting that same distance is circular.

TABLE IV
E3 SECOND-HOST TRANSFER ON ORB-SLAM3 (REAL C++ B0 VS. NATIVE ATLAS B2 VS. ZERO-LEAKAGE DUAL-ANCHORED B6)

| Sequence (GT) | Variant (Provenance) | Lost Frames | Submaps (Unmerged) | Global ATE (m) | Win. RMSE (m) | $G_{\mathrm{vs\,B2}}$ |
|---|---|---|---|---|---|---|
| (Leica)<br>$[120, 140]$ s<br>Severe ($g = 0.05$) | B0 Single-Map () | 174 | 1 (0) | 1.3755 | 1.1396 | — |
| | B2 Native Atlas () | 87 | 2 (0) | 2.1822 | 6.8941 | 0.0000 |
| | B5 Open-Loop () | 0 | 1 (0) | 6.5554 | — | −2.6306 |
| | **B6 Ours ()** | **0** | **1 (0)** | **0.9066** | **0.7366** | **+0.7674** |
| (Mocap)<br>$[25, 40]$ s Dark+Blur | B2 Native Atlas () | 80 | 2 (0) | 0.5902 | 1.4820 | 0.0000 |
| | **B6 Ours ()** | **0** | **1 (0)** | **0.2854** | **0.3120** | **+0.8637** |
| (Mocap)<br>(106.8 s Real Dark) | B0 Single-Map () | **904** | 1 (0) | 1.0186 | — | — |
| | B2 Native Atlas () | 181 | **5 (1 Unmerged)** | 0.4855 | — | 0.0000 |
| | B5 Concat () | 0 | 1 (0) | 0.6338 | — | −0.5975 |
| | **B6 Ours ()** | **0** | **1 (0)** | **0.2154** | — | **+1.0882** |
| **Held-Out:**<br>$[45, 65]$ s ($g = 0.05$) | B0 Single-Map () | **295** | 1 (0) | 1.3287 | 2.3376 | — |
| | B2 Native Atlas () | 168 | 2 (0) | 0.9739 | 1.8409 | 0.0000 |
| | **B6 Ours ()** | **0** | **1 (0)** | **0.6024** | **0.6057** | **+0.5262** |

### *D. E3 Second-Host Transfer to ORB-SLAM3 and Held-Out Blind Test (RQ5)*

Table IV and Fig. 6 evaluate second-host transfer on ORB-SLAM3 [3] using cross-scan point-to-plane holdout NIS

($\mathcal{S}_k \to \mathcal{S}_{k-5}$ transformed solely by the candidate trajectory's relative displacement $\Delta p_{\text{cand}}$) and paired Wilson sequential lower bounds. Under severe low-light ($g = 0.05$) on , real C++ single-map B0 () loses 174 frames (1.3755 m ATE), whereas native Atlas (B2) merges skewed transitional keyframes at 148 s and worsens global ATE to 2.1822 m (window RMSE 6.8941 m, cross-scan NIS $= 3.2854 > 2.5$). Our B6 repair achieves 0.9066 m global ATE and 0.7366 m window RMSE ($G_{\text{win}} = +0.9673$, $J_v^{\text{after}} = 116.51$, cross-scan NIS $= 2.1806 \leq 2.5$, paired Wilson LB $= 0.8162 \geq 0.55$). On the 106.8 s **real natural darkness** sequence (Mocap GT), single-map B0 permanently loses tracking at 48.13 s (1.0186 m ATE) and native Atlas leaves (74 frames) permanently unmerged (0.4855 m ATE, cross-scan NIS $= 1.4511$), whereas B6 achieves 0.2154 m ATE across all 1545 frames ($J_v^{\text{after}} = 142.34$, cross-scan NIS $= 1.1051 \leq 2.5$, paired Wilson LB $= 0.6919$). Finally, on the **held-out test sequence** (128.7 s, Leica GT, evaluated with strictly frozen development hyperparameters), MM-LINS misses the plastic-bag occlusion for a fourth consecutive sequence (=0, 4.0331 m $\to$ 0.2410 m B6, $G = +0.9864$), FAST-LIO2 recovers from 4.5275 m to $\mathbf{0.2796}$ m ($G = +0.9995$), LIO-SAM rejects the diverging B5 point filter (237.49 m) and recovers via B6 to $\mathbf{0.2430}$ m ($G = +0.9971$), and ORB-SLAM3 B6 reduces window RMSE from 1.8409 m (Atlas) to $\mathbf{0.6057}$ m ($G_{\text{win}} = +0.7782$, cross-scan NIS $= 2.1586$, Wilson LB $= 0.7902$).

On , the 130 s evaluation includes an injected 88% LiDAR occlusion during 50–70 s. FAST-LIO2 reports 1.2996 m after degradation and 0.1346 m for the repaired run, with 940 and 1,207 associated poses, respectively. MM-LINS reports 5.2635 and 0.3294 m on 237 and 236 poses. These differences describe the recorded runs; differing support prevents assigning the entire gain to the write decision. LIO-SAM's degraded trajectory reaches 333.29 m RMSE and the proposed offline bridge reaches 337.06 m. The bridge passes the boundary-jump criterion ($0.4042 \leq 0.50$ m) and is refused on the paired lower bound ($0.4626 < 0.55$). Its cross-scan score (4.4145) is not a calibrated criterion by itself: the clean run of the same sequence already scores 2.80–3.56 $> 2.5$ at the $\sigma = 0.25$ behind the archived number, and at $\sigma$ calibrated so that clean run scores 1 (0.42–0.47 m) the bridge scores 1.42. This calibration is an empirical normalisation fitted on the clean run, not a statistical noise model, and it is not yet validated on an independent interval. A statistic that rejects healthy data cannot refuse a candidate; the scale-free paired test is what refuses this one. The recorded decision is to wait.

In , naturally dark images are identical across bag variants; only the LiDAR stream receives the injected occlusion (15–35 s). On the shared post-occlusion interval (36.272–48.897 s), the FAST-LIO2 translation RMSE is 0.1053 m clean, 2.2022 m degraded and 2.0185 m for the constructed bridge. For LIO-SAM these values are 0.5076, 1.7124 and 2.1120 m. MM-LINS reaches 0.0828 m clean and 2.2811 m degraded, but its degraded keyframes begin at 34.57 s. Without a trusted pre-occlusion anchor, the bridge returns its input and repair is unevaluable. These outcomes distinguish a modest improvement, harm and a missing candidate. The evaluator shares a clean-run frame transform across each host's variants, whereas the stored protocol describes a separate pre-injection fit per run. We report the archived values pending reconciliation of this discrepancy. For the camera, the 0.3771 m position RMSE in the 13–40 s window comes from an offline bridge using healthy LiDAR motion. Its complete timestamps do not measure reacquisition by the visual tracker. On the gate's own statistic, the accepted candidate scores 0.65 with Wilson lower bound 0.924 at a noise scale calibrated on its clean run ($\sigma = 0.20$–$0.22$ m, where that clean run scores 0.63–0.77); at the simulator's 0.12 the same candidate scores 2.85 and would be refused.

### *E. Analytical 3D-Mesh Ray-Sampling Simulation (S1/S2, Provenance: ) and Runtime Overhead*

To evaluate true 3D volumetric map contamination against a finite bounded ground-truth surface mesh $\mathcal{G}_{\text{GT}}$ and execute controlled counterfactual gate tests for R4 (no loop revisit) versus R5 (delayed revisit to Zone A), we evaluate our online node ( in , with and validation reference points bound to the same transaction lifecycle as ) in an analytical 3D-mesh ray-sampling arena (, 40 m $\times$ 25 m, 5 seeds, provenance: , kept strictly distinct from real-host rosbag executions). Defining the true 3D trusted-map voxel contamination ratio at 0.10 m resolution as $\text{Contam}_{\text{3D}}(t) = |\mathcal{V}_{\text{trusted}}|^{-1} \sum_{v \in \mathcal{V}_{\text{trusted}}} \mathbb{I}(\min_{g \in \mathcal{G}_{\text{GT}}} \|v - g\|_2 > 0.15\,\text{m})$, unprotected B0 (43.25%), online projection B1 (43.21%), and same-data rollback B3 (43.25%) severely corrupt the map, while MM-LINS B2 ($\Delta t_{\text{margin}} = 0$ s) permanently writes $N_\kappa = 10$ confirmation-delay frames into the sleeping submap resulting in 17.79% 3D voxel contamination, and unvalidated single-modality B5 reaches 35.35% (Fig. 3b). Using every 5th scan as a strictly held-out scan set $\mathcal{S}_{\text{hold}}$ (Fig. 4), our 2.0 s delayed-commit buffer holds R4 in ($J_v^{\text{after}} = 0.79 < 25.0$), rejects narrow-FoV B5 ($J_v^{\text{after}} = 20.51$, rolling back its trial voxel writes via ), and also leaves B4 and B6 at $\text{Contam}_{\text{3D}} = 0.00\%$ — but here B6 admits nothing as well, so that figure says nothing about verified admission.

*a) Withdrawn result and simulator-validity check:* An earlier version of this manuscript reported R5 and B6 reaching in this arena, with cross-scan NIS of 1.4 and 0.8. That result is withdrawn. The validation reference had been assembled from the trusted map *plus the candidate's own non-holdout scans*, so the "independent" holdout was partly a comparison of the candidate against itself; once the reference is restricted to the pre-quarantine trusted map, every S1 candidate is rejected (R5 NIS $= 10.38$, B6 point-mixture 14.63, B6 pose-graph 10.38, all $> 2.5$) and the usable paired evidence halves from 20 to 10 frames of 20 keyframes.

Before reading that rejection as a finding about the method, we apply the same check to the ground truth. Placed at the true poses and scored against the same independent reference, the scans yield a mean holdout NIS of 18.9, 22.0, 12.6, 8.6 and 11.7 across the five scenarios — far above the 2.5 bound that any candidate must meet. (An earlier diagnostic reported medians of 21.2, 26.5, 13.8, 10.5 and 14.7 over all supported frames; that aggregation is not the gate's statistic and is retained only as history.) The analytic distance to the

arena surface gives a NIS of $0.0055$, so the geometry is self-consistent, but the discrepancy is not solely the normal estimator's: raising the reference budget from $1{,}500$ to $20{,}000$ moves the second scenario from $26.5$ to $14.8$, and larger budgets do not change the score, so density and neighbourhood structure are not excluded. A criterion the ground truth cannot satisfy cannot discriminate; we therefore mark the simulator as failing its validity check and cite no Mode 2 or Mode 3 verdict from S1/S2. The contamination ratios remain meaningful, since they measure what would be written rather than who decides; the gate verdicts do not.

This also exposed a limitation in the production path: truncates the reference cloud to $600$ points and the holdout scan to $300$ before fitting. Every holdout NIS reported in this paper, including those on recorded M2DGR sequences, was computed under that cap. We have not changed it, because doing so would alter results we have already reported; we flag it here so that it can be weighed.

On an Intel Core i7 CPU, our online health feature extraction, decoupled eigendecomposition, joint detector, and signature matrix router require $\mathbf{0.0599}\,\mathrm{ms}$ (P50) and $\mathbf{0.0852}\,\mathrm{ms}$ (P95) per window—only $\mathbf{0.0639}\%$ of a $10\,\mathrm{Hz}$ LiDAR frame budget—while adding the 48-cell spherical encoder and 2-layer causal relative-attention network takes $\mathbf{0.633}$–$\mathbf{0.639}\,\mathrm{ms}$ (P50). Delayed-commit queue management and SHA-256 checkpointing take $0.0034\,\mathrm{ms}$ (P50) and $0.4171\,\mathrm{ms}$ (P95) ($48.5\,\mathrm{KB}$ per snapshot). In a 10-scan window microbenchmark on the persistent $20\,\mathrm{cm}$ global voxel map of ($|M| = 594,979$ voxels, $|\Delta M_{\mathrm{dirty}}| = 82,219$ dirty voxels, $13.8\%$ of the map), full-dict deep-copy snapshot takes $1797.26\,\mathrm{ms}$, whereas our Copy-on-Write rolls back the $82,219$ dirty voxels in $\mathbf{99.20}\,\mathrm{ms}$ ($\mathbf{18.12}\times$ local rollback speedup) with exact bit-for-bit restoration. The 25k-point RANSAC mixture filter takes $30.82\,\mathrm{ms}$ (P50) and is invoked exclusively on quarantined background branches.

## VI. Limitations and Future Work

The information argument in Proposition 1 concerns a fixed linearization and fixed associations. When neither a revisit nor an auxiliary observation supplies usable geometry, the system can withhold writes; the experiments establish containment only in the conditions tested. Native host trajectories, replayed keyframe repairs and controlled mesh transactions measure different parts of the system. The manager can revoke represented keyframe and voxel edits; it does not demonstrate in-place rollback of the host's complete C++ filter or smoother. Point-cloud injection also leaves the optical response of a physical occluder unmeasured. The outstanding experiment is a single-host, same-input comparison of native mapping, quarantine alone, unconditional admission and verified admission. It must measure surface coverage alongside contamination, trajectory error on common support and unintended interventions on clean intervals before drawing a general map-recovery conclusion.

*a) Why point-level repair without verification is not merely inaccurate:* The near-wall hard-negative test deserves more weight than a single number. On un-injected , where nothing is wrong with the scene, an ungated geometry-only near-field filter strips real wall returns on $58/76$ ($76.3\%$) of near-wall scans; the NIS/$D_{\mathrm{LOSO}}$-gated router strips none on either $(0/76)$ or $(0/62)$. This failure differs in kind from the accuracy losses reported earlier, which at least occur on genuinely degraded input. Here the estimator is operating nominally and the repair still edits the map. A nearby wall and an occluder sheet are both planar surfaces one to three metres from the sensor, and nothing in the geometry of a single scan separates them: a filter that decides from range and planarity alone must either miss occluders or delete walls, and the choice of threshold only moves the error between those two outcomes. The consequences are also silent. After a false strip the map stays locally plausible, and the damage surfaces later as missing coverage, weaker constraints and eventual drift rather than as an immediately visible defect, so an operator has no reason to intervene. We therefore treat quarantine and verification as the enabling condition for point-level repair, not as an optional refinement of it, and every positive repair result in this paper is conditioned on the consistency gate being active. Conversely, the $76.3\%$ figure also bounds what our own results can claim: it shows the gate is necessary, not that it is sufficient, and the matched single-host comparison above is still required to establish how much geometry verified admission actually restores.

## VII. Conclusion

An estimator can resume tracking while the map still contains unsupported writes. Our keyframe transaction separates proposing a replacement from admitting it to the trusted map. Controlled trials show that delaying writes can protect occupied geometry. A recorded outdoor case shows that an apparently continuous bridge can nevertheless fail the geometric comparisons, while the common-window darkroom results show how candidate quality and even candidate availability depend on the host. The strongest remaining test is whether verified admission restores more independently measured surface than quarantine alone without sacrificing trajectory accuracy or intervening unnecessarily on clean input.

## References


[1] W. Xu, Y. Cai, D. He, J. Lin, and F. Zhang, "FAST-LIO2: Fast direct LiDAR-inertial odometry," *IEEE Transactions on Robotics*, vol. 38, no. 4, pp. 2053–2073, 2022.

[2] T. Shan, B. Englot, D. Meyers, W. Wang, C. Ratti, and D. Rus, "LIO-SAM: Tightly-coupled lidar inertial odometry via smoothing and mapping," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2020, pp. 5135–5142.

[3] C. Campos, R. Elvira, J. J. G. Rodríguez, J. M. M. Montiel, and J. D. Tardós, "ORB-SLAM3: An accurate open-source library for visual, visual–inertial, and multimap SLAM," *IEEE Transactions on Robotics*, vol. 37, no. 6, pp. 1874–1890, 2021.

[4] S. Zhao, H. Zhang, P. Wang, L. Nogueira, and S. Scherer, "Super odometry: IMU-centric LiDAR-visual-inertial estimator for challenging environments," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2021, pp. 8729–8736.

[5] Y. Ma, J. Xu, S. Yuan, T. Zhi, W. Yu, J. Zhou, and L. Xie, "MM-LINS: A multi-map LiDAR-inertial system for over-degenerate environments," *IEEE Transactions on Intelligent Vehicles*, vol. 10, no. 1, pp. 472–482, 2025.

[6] T. Tuna, J. Nubert, Y. Nava, S. Khattak, and M. Hutter, "X-ICP: Localizability-aware LiDAR registration for robust localization in extreme environments," *IEEE Transactions on Robotics*, vol. 40, pp. 452–471, 2024.
[7] J. Zhang, M. Kaess, and S. Singh, "On degeneracy of optimization-based state estimation problems," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2016, pp. 809–816.
[8] T. Shan, B. Englot, C. Ratti, and D. Rus, "LVI-SAM: Tightly-coupled lidar-visual-inertial odometry via smoothing and mapping," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2021, pp. 5692–5698.
[9] G. Kim and A. Kim, "Scan context: Egocentric spatial descriptor for place recognition within 3D point cloud map," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2018, pp. 4802–4809.
[10] H. Huang, W. Lin, and M. Liu, "Dual-SLAM: A framework for robust single-camera SLAM," *IEEE Transactions on Robotics*, vol. 36, no. 4, pp. 1142–1157, 2020.
[11] G. Kim and A. Kim, "LT-mapper: A modular framework for LiDAR-based lifelong mapping," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2022, pp. 7995–8002.
[12] N. Sünderhauf and P. Protzel, "Switchable constraints for robust pose graph SLAM," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2012, pp. 1879–1884.
[13] J. G. Mangelson, D. Dominic, R. M. Eustice, and R. Vasudevan, "Pairwise consistent measurement set maximization for robust multi-robot map merging," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2018, pp. 2916–2923.
[14] H. Yang, P. Antonante, V. Tzoumas, and L. Carlone, "Graduated non-convexity for robust spatial perception: From non-minimal solvers to global outlier rejection," *IEEE Robotics and Automation Letters*, vol. 5, no. 2, pp. 1127–1134, 2020.
[15] J. Yin, A. Li, T. Li, W. Yu, and D. Zou, "M2DGR: A multi-sensor and multi-scenario SLAM dataset for ground robots," *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 2266–2273, 2021.